
ML MODEL TO PREDICT CREDIT LIMIT

DATA

Credit card history records

Features
Age
Gender
Education_Level
Income_Category
Total_Amount_Checking
.....

Model



Target
Credit_Limit

PROCESS DATA

1. Choose meaningful features
2. Check continuous columns and categorical columns
3. Standardize continuous columns and impute them
4. OneHotEncode categorical columns and impute them
5. Construct the column transformer

CODE

```
target = credit['Credit_Limit']  
target = target.values.ravel()
```

```
categorical_columns = (X_train.dtypes == object)  
continuous_columns  = (X_train.dtypes != object)
```

```
con_pipe = Pipeline([('scalar', MaxAbsScaler()),  
                     ('imputer', SimpleImputer(missing_values=np.nan, strategy='median', add_indicator=True))  
                     ])  
  
cat_pipe = Pipeline([('ohe', OneHotEncoder(handle_unknown='ignore')),  
                     ('imputer', SimpleImputer(strategy='most_frequent', add_indicator=True))])  
  
preprocessing = ColumnTransformer([('categorical', cat_pipe, categorical_columns),  
                                   ('continuous', con_pipe, continuous_columns),  
                                   ])
```

TRAIN, TEST, PIPELINE

1. Split dataset into train set and test
 2. Use pca in the pipeline to focus on important features and improve generality
 3. Construct pipeline
-

CODE

```
X_train, X_test, y_train, y_test = train_test_split(data,  
                                                    target,  
                                                    test_size=0.2)
```

```
pipe = Pipeline([('preprocessing', preprocessing),  
                 ('pca', PCA(n_components=15)),  
                 ('clf', DummyEstimator())  
                ])
```

SEARCH BEST MODEL AND HYPERPARAMETERS

1. Construct search space by using 7 different models
2. Construct search algorithm
3. Search 5 times for the best model and hyper parameters

CODE

```
search_space = [{'clf': [RandomForestRegressor()],
    'clf__n_estimators': np.arange(100, 1000, 150), # decides how many trees
    'clf__max_features': ['log2', 'sqrt'], # decides how many features in each tree
    'clf__max_depth': np.arange(15, 25, 1), # decides how deep in each tree
    'clf__min_samples_leaf': np.arange(1, 10, 1), # decides how many samples at minimum in each leaf
    'clf__bootstrap': [True, False] # decides if using bootstrap technique
},

{'clf': [SVC()],
    'clf__C': np.logspace(0.1, 1000, 5), # decides the best C value
    'clf__gamma': np.logspace(0.0001, 1, 5), # decides the best gamma
    'clf__kernel': ['rbf', 'poly'], # decides kernel type
    'clf__class_weight': ['balanced', None] # decides the type of weighted class
},

{'clf': [RidgeCV()],
    'clf__normalize': [False, True], # decides if doing normalization
    'clf__alpha_per_target': [False, True] # decides if using alpha in per target
},

{'clf': [LassoCV()],
    'clf__eps': np.arange(0.0005, 0.01, 0.0005), # decides the best epsilon
    'clf__normalize': [False, True], # decides if doing normalization
    'clf__max_iter': np.arange(1000, 5000, 1000), # decides the maximum times of iteration
    'clf__n_alphas': np.arange(100, 500, 100) # decides the best n_alphas
},

{'clf': [BayesianRidge()],
    'clf__normalize': [False, True], # decides if doing normalization
    'clf__n_iter': np.arange(100, 1000, 100) # decides the times of iteration
},

{
    'clf': [HuberRegressor()],
    'clf__alpha': np.arange(0.0001, 0.001, 0.0001), # decides the best alpha
    'clf__max_iter': np.arange(100, 1000, 100), # decides the times of iteration
    'clf__epsilon': np.arange(1, 2, 0.1) # decides the best epsilon
},

{
    'clf': [ExtraTreesRegressor()],
    'clf__max_features': ['log2', 'sqrt'], # decides how many features in each tree
    'clf__max_depth': np.arange(15, 25, 1), # decides how deep in each tree
    'clf__n_estimators': np.arange(100, 1000, 150), # decides how many trees
    'clf__min_samples_leaf': np.arange(1, 10, 1), # decides how many samples at minimum in each leaf
    'clf__bootstrap': [True, False] # decides if using bootstrap technique
}
]

clf_algos_rand = RandomizedSearchCV(estimator=pipe,
    param_distributions=search_space,
    n_iter=50,
    cv=5,
    n_jobs=-1,
    verbose=10,
    scoring='neg_root_mean_squared_error')
```

```
for i in range(5):

    best_model = clf_algos_rand.fit(X_train, y_train)

    print(best_model.best_estimator_.get_params()['clf'], end='\n')
    print(best_model.best_score_, end='\n')
```

BEST MODEL AND HYPER PARAMETERS

1. Best model: RandomForestRegressor
2. Best hyperparameters
3. Final pipeline

CODE

```
params={'bootstrap': False,
        'ccp_alpha': 0.0,
        'criterion': 'mse',
        'max_depth': 20,
        'max_features': 'log2',
        'max_leaf_nodes': None,
        'max_samples': None,
        'min_impurity_decrease': 0.0,
        'min_impurity_split': None,
        'min_samples_leaf': 1,
        'min_samples_split': 2,
        'min_weight_fraction_leaf': 0.0,
        'n_estimators': 250,
        'n_jobs': None,
        'oob_score': False,
        'random_state': None,
        'verbose': 0,
        'warm_start': False}
```

```
pipe = Pipeline([('preprocessing', preprocessing),
                  ('pca', PCA(n_components=15)),
                  ('reg', RandomForestRegressor(**params))
                  ])
```

FIT MODEL AND PREDICT ON TEST SET

```
pipe.fit(X_train, y_train)
```

```
y_pred = pipe.predict(X_test)
```

```
mse = mean_squared_error(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)  
r_2 = r2_score(y_test, y_pred)
```

```
print('mse = {mse}\nmae = {mae}\nr2 = {r_2}'.format(mse=mse, mae=mae, r_2=r_2))
```

```
mse = 2575826.8301345236  
mae = 1044.6272158530253  
r2 = 0.9702055308283112
```

CONCLUSION

1. For this project, RandomForest performs best among RidgeCV, LassoCV, BayesianRidge, HuberRegressor, ExtraTreesRegressor, SVC
 2. $R^2 = 0.97$. Pretty good!
 3. RandomizedSearchCV has randomness
 4. PCA is helpful for increasing generality
 5. Do grid search later if possible
-

**THANK YOU FOR
WATCHING!**
